

Analysis of the Implementation of Cycle Generative Adversarial Network (CycleGAN)

Thadar Soe¹

¹Pusan National University

1. Introduction

CycleGAN, short for Cycle-Consistent Generative Adversarial Network, is a type of Generative Adversarial Network (GAN) designed to perform unpaired image-to-image translation.[4] Unlike traditional GANs that require paired datasets for training, CycleGAN is capable of learning mappings between two different image domains without needing paired examples. This makes it highly effective for tasks where paired data is unavailable or difficult to obtain.

CycleGAN uses two generators and two discriminators:

- Generator G_X : Transforms images from domain X to domain Y.
- Generator G_Y : Transforms images from domain Y to domain X.
- Discriminator D_X : Distinguishes between real images from domain X and generated images from domain X.
- Discriminator D_Y : Distinguishes between real images from domain Y and generated images from domain Y.

The generators and discriminators are trained simultaneously in an adversarial manner. The generators try to fool the discriminators by generating images that are indistinguishable from real images in their respective domains, while the discriminators aim to correctly classify whether an image is real or generated. [1]

2. Implementation

For this implementation, the MNIST dataset (domain X) and the SVHN dataset (domain Y) were used, where the MNIST Dataset is a set of handwritten digits, consisting of grayscale images with 1 channel, and the SVHN Dataset a set of street view house numbers, consisting of color images with 3 channels.

The CycleGAN model uses a combination of residual networks (ResNets) in the generator [2], and batch normalization [3] to stabilize training and improve performance. Three primary loss functions were used to train the model: Discriminator Loss (Loss_D), Generator Loss (Loss_G), and Cycle Consistency Loss (Loss_C). Each of these loss functions plays a crucial role in ensuring that the model learns to generate realistic images while maintaining the consistency of transformations between the two domains. Below is a detailed explanation of each of these loss components:

- Loss_D: The discriminator loss, representing how well the discriminator can distinguish between real and fake images. A lower value suggests the discriminator is performing well.
- Loss_G: The generator loss, representing how well the generator is creating realistic fake images. A lower value indicates the generator is improving.
- Loss_C: The cycle consistency loss, which ensures that the generator is capable of converting images from one domain to another and then back again to the original domain. Lower values indicate better cycle consistency.

3. Results

Over the course of training (as seen in Fig.2) from it [0100/4685] to it [4600/4685]), the losses are fluctuating but generally improving. The generator and discriminator are learning to generate more realistic images and distinguish them more accurately from real images.




Figure 1: Training Losses for CycleGAN




Figure 2: Generated images at Iterations 1000, 2000, 3000, 4000, and 4600 (top to bottom).

The discriminator loss initially fluctuates between values around 2.3 and decreases progressively to values closer to 1.2 by the later stages of training. This indicates that the discriminator becomes more proficient at distinguishing between real and fake images over time. However, periodic fluctuations occur due to the adversarial nature of the model, where the generator is constantly improving, forcing the discriminator to adapt and perform better.

Similarly, the generator loss begins relatively high and gradually decreases as the model learns to generate more realistic images. For instance, the generator loss starts at around 2.4 and decreases over time to values closer to 1.3 by iteration 4600. The decrease in generator loss implies that the generator is increasingly producing images that are difficult for the discriminator to differentiate from real images. This trend shows the generator's ability to learn the mapping between the two domains more effectively.

The cycle consistency loss shows a generally low and stable trend, particularly towards the later iterations, indicating that the model is maintaining the original structure of images after translating them from one domain to the other and then back again. This is a crucial aspect of the CycleGAN model, as the cycle consistency ensures that the transformations made by the generator preserve the key features of the images.

Alongside the numerical improvements in the loss functions, the qualitative visual results from the generated images also show significant progress. Early generated images may still exhibit noticeable distortions, but as training progresses, these artifacts decrease, and the generated images become more indistinguishable from real images in both domains. For example, at iterations 1000, 2000, and 3000, the images might still contain visible artifacts and noise, but by iteration 4600, the generated images are much sharper and more coherent. These visual improvements reflect the effectiveness of the adversarial training and the cycle consistency loss in guiding the generators to produce realistic outputs.

4. Citations

- [1] Ian Goodfellow, Jean Pouget-Abadie, Mehdi Mirza, Bing Xu, David Warde-Farley, Sherjil Ozair, Aaron Courville, and Yoshua Bengio. Generative adversarial nets. In *Advances in Neural Information Processing Systems*, volume 27, 2014. URL <https://arxiv.org/abs/1406.2661>.
- [2] Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. Deep residual learning for image recognition. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pages 770–778, 2016. URL <https://arxiv.org/abs/1512.03385>.
- [3] Sergey Ioffe and Christian Szegedy. Batch normalization: Accelerating deep network training by reducing internal covariate shift. In *International conference on machine learning*, pages 448–456, 2015. URL <https://arxiv.org/abs/1502.03167>.
- [4] Jun-Yan Zhu, Taesung Park, Phillip Isola, and Alexei A Efros. Unpaired image-to-image translation using cycle-consistent adversarial networks. *arXiv preprint arXiv:1703.10593*, 2017. URL <https://arxiv.org/abs/1703.10593>.